

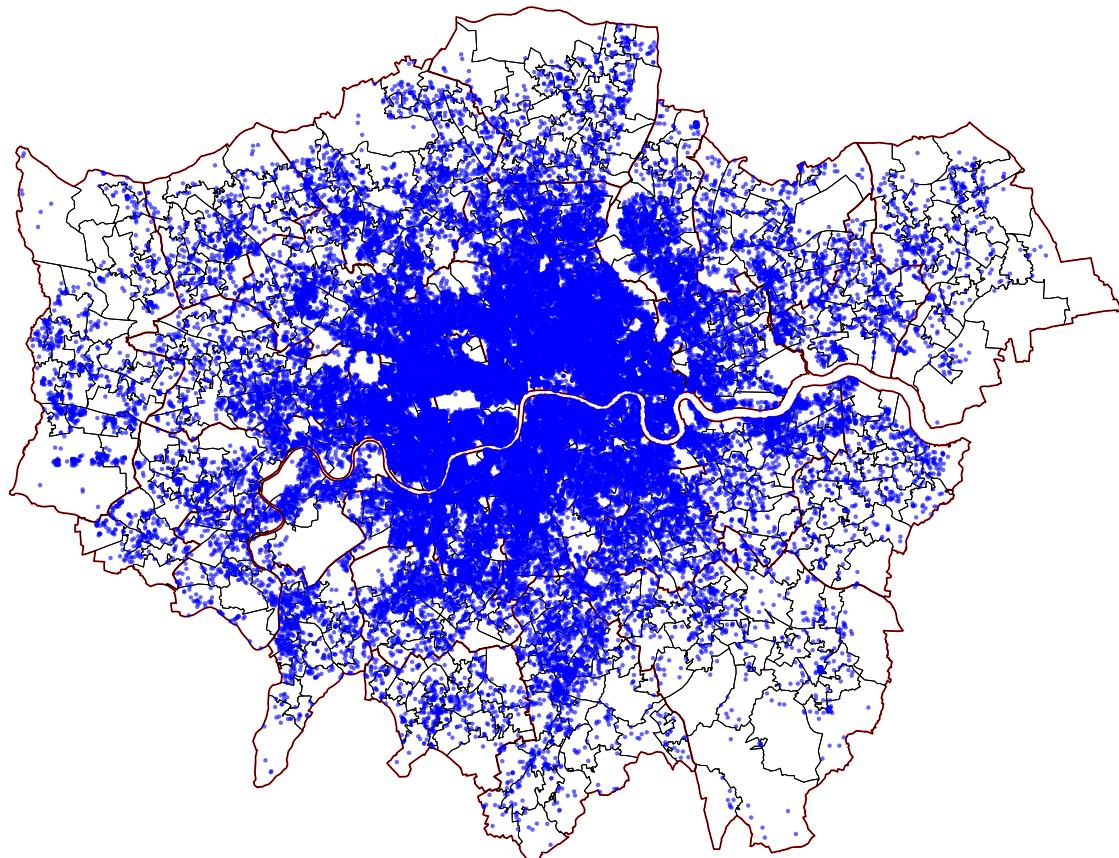
Spatial Analysis of Professional Airbnb Activity in London

Christian Mulder


2025-11-01

To begin with, I will start by reading in all the necessary files, which are the listings CSV, the borough boundary data, and then the MSOA data.

All Listings over Borough (red) and MSOA (black) Boundaries



Professional Hosts (>=2 listings AND availability >=90 days) over Borough (red) and MSOA (black) Boundaries



The reduced dataset for professional landlords is 41664

In order to further analyse, I will take these following steps: - Counts and distributions (MSOA/borough, normal/log). - Normality checks (Shapiro–Wilk). - Gini coefficient (MSOA and borough) → here it acts as a descriptive inequality measure alongside distributions. - Density visualization (KDE attempt, maps). - Global spatial autocorrelation (Moran's I). - Neighbor definition comparison (KNN vs Distance Band). - Spatial Error Model. - *Local clustering (Getis–Ord Gi)**. - Return to inequality in discussion → you can revisit the Gini in the conclusion, tying it to spatial clustering.

```
Index(['id', 'listing_url', 'last_scraped', 'name', 'description', 'hos  
'host_name', 'host_since', 'host_location', 'host_is_superhost',  
'host_listings_count', 'host_total_listings_count',  
'host_verifications', 'latitude', 'longitude', 'property_type',  
'room_type', 'accommodates', 'bathrooms', 'bathrooms_text', 'bed
```

```
'beds', 'amenities', 'price', 'minimum_nights', 'maximum_nights',
'availability_365', 'number_of_reviews', 'first_review', 'last_review',
'review_scores_rating', 'reviews_per_month'],
dtype='object')
```


Counts and Distributions

First, the following two maps will be the distribution of counts (tallied per area) for the Borough and MSOA data.

#####MSOA:




Professional Host Listings per MSOA (boroughs in red)



Professional Host Listings per Borough

Borough:




Distribution table:

Just to test whether this is a bad skew (although I would say the diagram is pretty obvious...), I will do a Shapiro-Wilk p-value test:

Shapiro-Wilk p-value: $4.1282625722644046e-47$

Because from this bar chart of the normal distribution, which is very extremely positively skewed, I want to try creating a log-scaled histogram to see if that changes anything.




And again let's test the Shapiro-Wilk p-value:


Shapiro-Wilk p-value (log counts): 0.0002740993476703627

The Shapiro–Wilk test on the raw MSOA counts returned a p□value of 2.3×10^{-46} , providing overwhelming evidence against normality. This confirms what the histogram suggested: the distribution of multi□host listings is extremely skewed. After applying a log transformation, the p□value increased to 0.0300. Although this is still below the conventional 0.05 threshold (meaning the log□transformed data cannot be considered perfectly normal), the improvement is substantial. The log scale reduces skewness and brings the distribution closer to normality, making it more suitable for visualization and interpretation, especially when comparing areas with smaller counts.

To better visualise the distribution, the following map applies a log scale transformation. This adjustment reduces the dominance of areas with exceptionally high counts (such as Westminster), allowing spatial patterns in other boroughs to remain visible and interpretable.


Professional Host Listings per MSOA (log scale)






Gini I will now do a quick Gini test, which measures how unevenly multi-host listings are distributed across areas. A higher Gini value indicates that activity is concentrated in fewer places, while a lower value suggests a more even spread.

The Gini coefficient ranges from 0 (perfect equality) to 1 (maximum inequality), with values above about 0.5 generally considered high.




The Gini coefficient of 0.503 at the borough scale indicates a moderate degree of inequality in the distribution of multi-host listings across London. Some boroughs host substantially more multi-host activity than others, but the overall pattern is less extreme than at finer spatial resolutions.



At the MSOA scale, the Gini coefficient rises to 0.627, revealing a much higher degree of spatial inequality. This shows that inequality is not only present between boroughs, but also concentrated within them: a relatively small number of MSOAs account for a disproportionately large share of multi-host listings, while the majority contain very few or none. In other words, hotspots of multi-host activity emerge at the neighbourhood level, highlighting that regulation or policy responses cannot rely solely on borough-wide averages but must account for highly localized concentrations.

———— ## Autocorrelation & Point Pattern Analysis

KDE Kernel Density Estimation (KDE) allows you to move beyond raw counts by creating a smooth surface that highlights where multi-host listings are most concentrated. This helps reveal localized hotspots and spatial patterns that might be obscured when only looking at aggregated counts.




Now the following map has decided on fixed boundaries:

So I think this simpler map might be better for displaying a simplified version of the KDE considering who we are presenting it to. However, I will need to continue working on this, as the heat map does not seem to correspond somehow with the actual results, so maybe I have made it too smooth and diluted in my attempt to accurately match it to the boundary lines. So it will need some more work.

Now because I find that both these maps look like they're a little bit off when I'm comparing them to the count (although that is possible in KDE because it is different) I will definitely continue from here. For the remaining autocorrelation I will first do Moran's I:

Moran's I Test (Local and Global Outcomes)

Moran's I: 0.6356
 p-value (permutation): 0.0010
 z-score (normal approx): 42.039
 MSOAs with zero neighbors (should be 0): 0




Moran's I of 0.6465 with a p-value of 0.001 indicates strong, statistically significant spatial clustering of multi-host listings across London MSOAs. High-count areas are surrounded by other high-count areas, confirming that the distribution is far from random. This suggests concentrated zones of multi-host activity that may warrant targeted policy attention.

Having confirmed spatial autocorrelation with Moran's I, I will now test the spatial lag and error models to assess whether the variation in multi-host density is better explained by neighboring effects or unobserved spatial structure in the census-like boundary data.

Spatial Lag Model

Spatial Lag Model (counts only)
 ρ (spatial lag coefficient): 0.9214
 Pseudo R²: 0.7124
 AIC: 10665.28
 Intercept: 4.2620



While I ran the spatial lag model to explore potential spillover effects between neighboring MSOAs, I don't believe it adds meaningful insight in this context. Just because two areas are adjacent doesn't imply that multi-host listing counts in one directly influence the other — these are not inherently relational phenomena. The clustering observed is more likely due to shared local conditions than spatial dependence, so including the lag model may misrepresent the underlying dynamics.

Since adjacency alone doesn't imply causal influence between MSOAs, I turn instead to a spatial error model to capture unobserved spatial clustering in the residuals.

Spatial Error Model

REGRESSION RESULTS

SUMMARY OF OUTPUT: GM SPATIALLY WEIGHTED LEAST SQUARES

Data set	:	unknown	
Weights matrix	:	unknown	
Dependent Variable	:	prof_count	Number of Observations
Mean dependent var	:	43.5353	Number of Variables
S.D. dependent var	:	79.2715	Degrees of Freedom

Pseudo R-squared : 0.1148

Statistic	Variable	Coefficient	Std.Error	z-
		Probability		
	CONSTANT	-86.67725	42.49712	-
2.03960	0.04139			
	USUALRES	-0.01970	0.00412	-
4.78764	0.00000			
	POPDEN	-0.17658	0.04421	-
3.99438	0.00006			
	HHOLDS	0.06787	0.01021	6.64758
	AVHHOLDSDZ	36.67220	14.71466	2.49222
	lambda	0.93857		

===== END OF REPORT =====


The model shows that while household variables like number of households and average household size are positively associated with multi-host listings, population and density are negatively associated. However, the most striking result is the very high and significant spatial error coefficient ($\lambda \approx 0.93$), which indicates that even after accounting for these demographic factors, there remains strong spatial autocorrelation in the residuals. In other words, clustering of multi-host activity is not fully explained by population or housing characteristics alone, but is (obviously) impacted by another variable. If we could find some data on tourism, housing cost, or anything else like this that could definitely further expand this.

However, crucially this above output is not very presentable to a mayor. So while we can maybe discuss it, saying our own research showed that these variables used show that there are other factors that need to be accounted for, I think it would be better not to show the output as it might not be visually pleasing (like even I don't wanna look at it realistically).

Getis-Ord Gi Next, I will use the Getis-Ord Gi statistic to identify local hotspots and coldspots of multi-host activity.

	MSOA11CD	count	GiZ	GiP	GiSig
0	E02000001	392	3.170811	0.002	Hotspot
1	E02000001	392	2.106223	0.010	Hotspot
2	E02000002	1	-0.373221	0.096	Not Significant
3	E02000003	15	-0.362928	0.116	Not Significant
4	E02000004	4	-0.320205	0.194	Not Significant

Getis-Ord Gi* Hotspots of Professional Host Properties



So because the MSOA data for some reason doesn't recognise that the polygons are connected to one another (so it thinks they are separate) I had to define the number of neighbours (which I set to 5 so it isn't too many and over the top of a strange pattern) but this does mean this map is not

entirely accurate. It is still interesting though.

Now I will do the same with the boroughs, which doesn't have that problem of neighbouring. I think this could be good to show initially (so it is clear where the hotspots and relevant boroughs are as I think this is more important for policy than MSOAs)

Getis-Ord Gi* Hotspots of Professional Host Properties (Normalized by Area)



This map uses the standard Getis-Ord Gi* method with Queen contiguity and a 95% confidence threshold, which is widely accepted in spatial analysis. The input values are normalized by borough area (per km²), ensuring fair comparison across differently sized regions. While this conservative approach may underdetect coldspots, it ensures statistical rigor and avoids overinterpreting weak patterns. For exploratory purposes, a relaxed threshold (e.g. 90%) can be used in parallel to surface borderline cases.